

Hypergraph attention reduces computational requirements.

Sparse quantization preserves latency on consumer hardware.

Mixed precision training is a tradeoff of speed and accuracy.

Our evaluation also examines throughput at small batch sizes.