

Approximate attention accelerates long-context inference.

Low-rank pruning slashes parameter count without retraining.

Drastic compression remains brittle against noisy inputs.